Decomposing Alpha: A Deep Dive into Wavelet-Based Factor Timing

Part I: Theoretical Foundations of Wavelet Analysis in Finance

1.1 The Challenge of Non-Stationarity in Financial Time Series

Financial time series—such as the returns of assets, factors, or portfolios—exhibit statistical properties that are fundamentally distinct from the signals commonly encountered in classical physics and engineering. The most critical of these properties is non-stationarity: the mean, variance, and covariance of financial data are not constant but evolve over time.[1] This time-varying nature manifests as volatility clustering, where periods of high turbulence are followed by periods of relative calm; structural breaks, where the underlying data-generating process shifts abruptly due to economic shocks or policy changes; and complex trends that are not simple, deterministic functions of time.[3]

This inherent non-stationarity poses a significant challenge to traditional time series analysis techniques. Methods rooted in the frequency domain, most notably the Fourier Transform, are particularly ill-suited for this environment. The Fourier Transform decomposes a signal into its constituent sine and cosine waves, providing exceptional resolution in the frequency domain. It can precisely identify *which* frequencies are present within a signal's entire duration. However, this precision comes at the cost of complete loss of information in the time domain; it cannot specify *when* those frequencies occurred.[4] This is a fatal flaw for financial analysis, as the assumption that market activity is stable across an entire period is demonstrably false.[3] For instance, a Fourier analysis might reveal a high-frequency component in a factor's return series, but it cannot distinguish between a brief, violent market crash and persistent, low-level noise spread over decades.

The assumption of stationarity is a foundational pillar for many widely used econometric models, such as the Autoregressive Integrated Moving Average (ARIMA) family. When this assumption is violated, models can produce spurious results, indicating relationships between variables where none exist, leading to unreliable forecasts and flawed conclusions.[1] While techniques like differencing can induce stationarity in some cases (e.g., transforming a random walk into a stationary process), they can also filter out important long-term trend information.[1] Modern research continues to grapple with the complexities of non-stationarity, recognizing its potential to obscure crucial long-term relationships while creating misleading short-term correlations.[6] It is within this context of non-stationary, event-driven, and multi-faceted financial data that more advanced signal processing tools become not just advantageous, but necessary.

1.2 The Wavelet Transform: A Multi-Resolution Lens for Financial Signals

The wavelet transform emerges as a powerful alternative, specifically designed to overcome the limitations of traditional methods when analyzing non-stationary signals. Its principal advantage lies in its ability to provide simultaneous localization in both the time and frequency domains, offering a time-frequency representation of a signal.[3]

Core Advantage: Time-Frequency Localization

Unlike the infinitely long sine and cosine waves used by Fourier analysis, the basis functions of the wavelet transform are small, localized waves called "wavelets." These functions are generated from a single prototype function, the "mother wavelet," through two fundamental operations: scaling (dilation or compression) and shifting (translation).[5] Scaling the mother wavelet allows for the analysis of different frequencies; a dilated (stretched) wavelet is sensitive to low-frequency, long-term phenomena, while a compressed wavelet captures high-frequency, short-term events. Shifting the scaled wavelet along the time axis allows for the analysis of these frequency components at specific points in time. The result is a two-dimensional representation of the signal, mapping its frequency content as it evolves through time. This makes it possible to identify transient events, such as market shocks, and analyze their frequency characteristics precisely when they occur.[5]

Multi-Resolution Analysis (MRA)

A key framework for applying wavelets is Multi-Resolution Analysis (MRA). MRA formalizes the idea of decomposing a signal into different "levels" or "scales" of resolution.[10] At each level, the signal is split into a low-frequency component, known as the "approximation," and a high-frequency component, known as the "detail." The approximation from one level becomes the input for the next, creating a hierarchical decomposition that separates the signal's long-term, underlying trend from its finer, short-term fluctuations.[13] This structure aligns perfectly with the multi-scale nature of financial markets, where different economic forces and market participants operate on distinct time horizons—from high-frequency traders reacting in seconds to institutional investors with multi-year outlooks.[7]

Types of Wavelet Transforms

The theoretical framework of wavelet analysis is implemented through several types of transforms, each with specific properties suited to different applications.

- **Continuous Wavelet Transform (CWT):** The CWT computes the wavelet coefficients at every possible scale and time shift, providing a highly detailed and redundant time-frequency map of the signal.[3] This redundancy makes it an excellent tool for visualization, particularly through scalograms (heatmaps of wavelet power), and for detailed analysis of the time-frequency signature of specific, isolated events. However, the high computational cost and redundancy make it less practical for applications requiring efficient signal reconstruction or feature extraction from large financial datasets.[3]
- **Discrete Wavelet Transform (DWT):** The DWT is a more computationally efficient implementation that evaluates the transform at a discrete set of scales and time shifts, typically on a dyadic grid (powers of two).[3] This process yields a non-redundant, compact representation of the signal, making the DWT the standard choice for tasks such as denoising, data compression, and feature extraction for machine learning models. Its main drawback is its
  shift-variance: a small shift in the input signal's starting point can lead to a significant change in the distribution of energy across wavelet coefficients at different scales.
- **Maximal Overlap DWT (MODWT):** Also known as the undecimated or stationary wavelet transform, the MODWT is a variant of the DWT that overcomes the issue of shift-variance.[3] It achieves this by not downsampling the coefficients at each level of decomposition, resulting in a redundant but shift-invariant transform. This property is highly desirable for financial time series analysis, as the results of the analysis should not depend on the arbitrary start date of the data. A particularly powerful feature of the MODWT is its ability to partition the variance of the original time series on a scale-by-scale basis. This allows for a rigorous analysis of how much of the signal's total volatility is attributable to phenomena at different time horizons (e.g., short-term cycles vs. long-term trends).[11]

1.3 A Comparative Analysis of Signal Decomposition Techniques

The decision to employ wavelet analysis for decomposing a factor return series is a deliberate one, reflecting a trade-off between model flexibility and underlying assumptions. To fully appreciate this choice, it is instructive to compare the wavelet approach to two other prominent methods used for trend-cycle decomposition in economics and finance: the Hodrick-Prescott (HP) filter and the Kalman filter.

Hodrick-Prescott (HP) Filter

The HP filter is a non-parametric data-smoothing tool widely used in macroeconomics to separate a time series into a trend component and a cyclical component.[14] It achieves this by solving a convex optimization problem that minimizes the sum of the squared deviations of the series from its trend, subject to a penalty on the variation in the trend's growth rate.[16] This penalty is controlled by a single smoothing parameter,
λ.

- **Underlying Assumption:** The core assumption of the HP filter is that the true underlying trend of a series is a smoothly varying component. The filter is formally optimal for time series that are integrated of order two, i.e., I(2) processes.[15]
- **Limitations:** The HP filter has several well-documented drawbacks for financial applications. First, it is a two-sided filter, meaning that the estimation of the trend at time t uses data from both before and after t. This non-causal nature makes it unsuitable for any real-time forecasting or trading application, as it relies on future information.[15] Second, the filter is known to perform poorly at the endpoints of a series, where the trend estimates can be highly unreliable and subject to significant revision as new data becomes available.[16] Finally, it has been shown that the HP filter can induce spurious cyclicality and dynamic relationships in filtered data that are not present in the original data-generating process.[15]

Kalman Filter

The Kalman filter is a sophisticated recursive algorithm that provides an optimal estimate of the hidden state of a dynamic system from a sequence of noisy and incomplete measurements.[17] It operates through a two-step process: a
prediction step, where it projects the current state and error covariance estimates forward to obtain an *a priori* estimate for the next time step, and an *update* step, where it refines this estimate using the next actual measurement.[17]

- **Underlying Assumption:** The primary requirement for the Kalman filter is that the system dynamics can be accurately described by a linear state-space model with Gaussian noise terms.[18] This is a strong and often restrictive assumption for complex, non-linear financial processes.
- **Application:** The Kalman filter is extremely versatile and can be used for filtering (estimating the current state), smoothing (estimating past states), and forecasting (estimating future states). Many standard time series models, including ARMA processes, can be cast into a state-space representation, allowing them to be estimated via the Kalman filter.[18] As a recursive, one-sided filter, it is inherently causal and well-suited for real-time applications.

The choice among these methods reveals a fundamental tension in quantitative modeling. The HP filter imposes a rigid, *a priori* belief about the smoothness of the trend, which may be inappropriate for financial markets characterized by abrupt regime shifts. The Kalman filter demands the specification of a precise parametric model for the underlying data-generating process; if this model is misspecified, the resulting decomposition will be flawed.

Wavelet analysis occupies a middle ground. It is a non-parametric method that, unlike the HP filter, does not impose a global smoothness constraint. Unlike the Kalman filter, it does not require strong assumptions about the functional form of the underlying process.[3] It is a more flexible, data-driven approach that allows the properties of the signal itself to determine the decomposition. This flexibility makes it an ideal tool for the exploratory analysis proposed in this project, where the goal is to extract a signal without imposing strong prior beliefs.

However, this very flexibility introduces a new set of challenges. The numerous choices involved in a wavelet analysis—the mother wavelet, the decomposition level, the denoising method—create "researcher degrees of freedom" that, if not handled with discipline, can lead to data snooping and overfitting. This critical issue will be a central theme in the subsequent analysis of the project's robustness.

Table 1: Comparison of Signal Decomposition Methods

Feature	Wavelet Analysis	Kalman Filter	Hodrick-Prescott (HP) Filter
Core Principle	Multi-resolution decomposition using scaled and shifted basis functions.	Recursive Bayesian estimation of the hidden state of a linear dynamic system.	Penalized least squares minimization to separate a smooth trend from a cyclical component.
Key Assumptions	No strong assumption on the data-generating process; assumes signal has features that match the chosen wavelet.	System can be described by a linear state-space model with Gaussian noise.	The underlying trend is a smoothly varying component; optimal for I(2) processes.
Causality	Can be implemented causally (using only past data), making it suitable for real-time applications.	Inherently causal and recursive; ideal for real-time filtering and forecasting.	Non-causal (two-sided filter); uses future data to determine the current trend, making it unsuitable for forecasting.
Strengths for Financial Data	Excellent time-frequency localization for analyzing non-stationary signals and transient events	Optimal for systems with a known linear structure. Can handle missing observations and incorporate multiple data sources.	Simple to implement and widely understood in macroeconomics.

	(e.g., shocks). Flexible and data-driven.		
Weaknesses/Risks	Many researcher degrees of freedom (choice of wavelet, level) create a high risk of overfitting and data snooping. Boundary effects can distort results.	Performance is highly sensitive to the accuracy of the state-space model specification. Assumes linearity and Gaussianity, which may not hold.	Can induce spurious cycles. Performs poorly at the endpoints of a series. The choice of the smoothing parameter (λ) is often arbitrary.

Part II: A Practical Guide to Project Implementation

This section provides a detailed, step-by-step methodology for executing the proposed two-week research project. It covers the entire workflow, from data acquisition and factor construction to the application of the wavelet transform and the final backtesting of the timing strategy. Each step is accompanied by a discussion of the critical methodological choices and practical considerations required for a rigorous implementation.

2.1 Step 1: Constructing the Momentum Factor

The foundation of this project is a reliable, well-constructed time series of momentum factor returns. This process is far from trivial and involves careful data acquisition, linking, cleaning, and a precise implementation of the factor definition.

Data Acquisition and Integrity

- **Primary Data Sources:** For a comprehensive analysis of the US stock market, two primary databases are indispensable: the Center for Research in Security Prices (CRSP) database, which provides high-quality stock price, return, and volume data, and a source for fundamental company accounting data, such as S&P's Compustat or FactSet.[20] The combination of CRSP's market data with Compustat's fundamental data allows for the construction of a wide array of factors.
- **Database Linking:** A significant technical challenge is linking the securities in CRSP with the companies in Compustat, as they use different primary identifiers (PERMNO and PERMCO in CRSP, GVKEY in Compustat). This process is complicated by corporate actions like mergers, acquisitions, and name changes, which can alter identifiers over

time.[23] For this project, the standard solution is to use the CRSP/Compustat Merged (CCM) database, typically accessed via Wharton Research Data Services (WRDS), which provides a well-maintained historical link between these identifiers.[23]

- **Data Quality Challenges:** It is a critical error to assume that commercial databases are free from errors or biases.[24] Compustat, while an industry standard, presents several known challenges that can impact research outcomes:
  1. **Standardization Effects:** Compustat applies a proprietary set of rules to standardize raw accounting data from company filings to make them comparable across firms. However, this process can introduce significant discrepancies when compared to the "as-filed" data available in machine-readable XBRL format from the SEC. Research has shown that these discrepancies can be large enough to alter the conclusions of asset pricing tests; for example, one study found that the well-known accruals anomaly was statistically significant when using as-filed data but disappeared when using Compustat data for the same period.[26]
  2. **Vendor Divergence:** Different data vendors (e.g., Compustat, FactSet, Refinitiv) employ different standardization methodologies. This can lead to materially different values for the same accounting item, resulting in different factor scores, different portfolio constituents, and ultimately, different backtested returns.[25]
  3. **Limited Scope:** Compustat data primarily covers publicly traded firms, which means that measures of industry concentration or competition derived solely from this data can be skewed by the omission of large private firms.[29]

Data Cleaning and Preparation Protocol

Before any factor can be constructed, the raw data must undergo a rigorous cleaning and preparation process to ensure its integrity. This protocol is a critical and often time-consuming part of quantitative research.[30] A robust protocol includes the following steps:

- **Universe Definition:** Restrict the sample to ordinary common shares (CRSP SHRCD = 10 or 11) trading on major US exchanges (NYSE, AMEX, NASDAQ; CRSP EXCHCD = 1, 2, or 3). This removes instruments like ADRs, REITs, and closed-end funds that have different characteristics.
- **Handling Missing Data:** Ensure that firms have the necessary data for factor construction (e.g., sufficient price history for momentum). Missing returns or prices must be handled consistently. CRSP uses special codes (e.g., -99.0) for missing data that must be correctly interpreted.[33]
- **Delisting Returns:** When a stock is delisted, CRSP provides a delisting return. It is crucial to incorporate these returns to avoid survivorship bias, as delistings are often associated with poor performance.
- **Outlier Treatment:** Financial databases can contain data entry errors that manifest as extreme outliers. A common practice is to winsorize or truncate variables at extreme percentiles (e.g., the 1st and 99th percentiles) each cross-sectionally to mitigate the

influence of these erroneous data points.[22] This should be done with care, as some extreme values may reflect genuine economic events rather than errors.[34]

Factor Definition: 12-Month Momentum (UMD)

This project will construct the standard 12-month momentum factor, often referred to as UMD (Up-Minus-Down), following the canonical methodology established by Jegadeesh and Titman (1993) and Fama and French.[20]

- **Universe and Eligibility:** At the end of each month t−1, the universe consists of all eligible stocks from the cleaned CRSP database. To be included, a stock must have a valid price at the end of month t−13 and a valid return for month t−2 to ensure the momentum period is calculable.[33]
- **Momentum Signal Calculation:** For each stock, the momentum signal is its cumulative return over the 11-month period from the end of month t−12 to the end of month t−2. The most recent month (t−1) is skipped to separate the momentum effect from short-term reversal patterns commonly observed in stock returns.[33]
- **Portfolio Formation:** At the end of each month t−1, all eligible stocks are ranked based on their momentum signal. To control for the microstructure effects of smaller stocks, the ranking breakpoints (e.g., for deciles or quintiles) are determined using only NYSE-listed stocks. The 30th and 70th NYSE percentiles are common breakpoints.[33] All stocks (NYSE, AMEX, and NASDAQ) are then allocated to portfolios based on these NYSE breakpoints.
- **Factor Return Calculation:** The momentum factor is a zero-investment, long-short portfolio. The monthly return of the UMD factor is calculated as the average return of the stocks in the top-performing portfolio (the "winners") minus the average return of the stocks in the bottom-performing portfolio (the "losers") for the subsequent month, t. The portfolios can be either equal-weighted or value-weighted (using market capitalization from the end of month t−1). The Fama-French methodology uses a 2x3 bivariate sort on size and prior return to construct six value-weighted portfolios, and the momentum factor is then calculated as Mom=21(SmallHigh+BigHigh)−21(SmallLow+BigLow).[33] This project will primarily use an equal-weighted decile-spread portfolio for simplicity, but the Fama-French construction serves as an important robustness check. This procedure, repeated monthly, generates the raw time series of momentum factor returns that will be the input for the wavelet analysis.

2.2 Step 2: Applying the Wavelet Transform

With the momentum factor time series constructed, the next step is to apply a wavelet transform to decompose it into its constituent time-frequency components. This involves several critical design choices that will significantly influence the final result.

Practical Implementation in Python

The PyWavelets library is the standard open-source tool for wavelet analysis in Python.[12] The primary function for this project is
pywt.wavedec(), which performs a multilevel Discrete Wavelet Transform (DWT). The function takes the signal, a chosen mother wavelet, and the desired decomposition level as inputs and returns a list of coefficient arrays.

Critical Design Choice 1: The Mother Wavelet

The choice of the mother wavelet is arguably the most important decision in a wavelet analysis, as the properties of the wavelet determine what features can be extracted from the signal.[3] The ideal wavelet should have properties that match the characteristics of the signal being analyzed. For financial time series, properties like smoothness and symmetry are often desirable.

- **Haar:** The simplest wavelet, discontinuous and resembling a square wave. Its strength is its perfect localization in time, making it ideal for detecting abrupt, step-like changes or jumps.[3] However, its lack of smoothness makes it a poor choice for representing the smoother, cyclical components often sought in factor returns.
- **Daubechies (dbN):** A family of orthogonal wavelets with compact support. The integer 'N' refers to the order; higher orders are smoother but have wider support (meaning they are less localized in time) and are more computationally intensive. A db4 or db6 is often a good starting point.[12]
- **Symlets (symN) and Coiflets (coifN):** These are near-symmetric modifications of the Daubechies wavelets. Symmetry is an important property because it ensures that the wavelet transform does not introduce phase distortion, meaning that features in the decomposed signal align correctly in time with the features of the original signal.[12] For this project, a mid-order Symlet, such as
sym8, represents a sound default choice, offering a good balance between smoothness, compact support, and near-perfect symmetry.

Table 2: Characteristics of Common Wavelet Families

Wavelet Family	Key Properties	Best Suited For	Potential Drawbacks
Haar	Discontinuous, orthogonal, symmetric, compact support.	Detecting sharp, step-like discontinuities and jumps.	Poor frequency resolution; not smooth, leading to aliasing.
Daubechies (dbN)	Orthogonal, asymmetric, compact	General-purpose signal analysis,	Asymmetry can cause phase distortion in

	support. Smoothness increases with order N.	compression. Good energy compaction.	signal reconstruction.
Symlets (symN)	Near-symmetric, orthogonal, compact support. A modification of Daubechies.	Applications where phase distortion is a concern. Good balance of properties.	Not perfectly symmetric (except for Haar, which is sym1).
Coiflets (coifN)	Near-symmetric, orthogonal, compact support. Vanishing moments for both wavelet and scaling function.	Similar to Symlets, useful for reconstruction tasks.	No significant advantage over Symlets for many applications.
Morlet	Non-orthogonal, complex-valued, symmetric. Good time and frequency resolution.	Continuous Wavelet Transform (CWT) for detailed time-frequency analysis and visualization.	Not suitable for standard DWT-based reconstruction; computationally intensive.

Critical Design Choice 2: The Decomposition Level

The decomposition level determines the number of scales into which the signal is broken down, effectively defining the frequency bands of the analysis. Choosing an appropriate level is a trade-off: too few levels may fail to isolate the high-frequency noise from the underlying signal, while too many levels may erroneously filter out meaningful short-to-medium term cyclical behavior.[36]

There is no universally accepted rule for selecting the optimal level. A common heuristic for a time series of length N is to choose a maximum level of $L=\lfloor \log2(N) \rfloor$.[37] For a typical factor return series spanning several decades (e.g., 50 years of monthly data gives $N=600$), this would suggest a maximum level of 9.

A more data-driven approach involves examining the distribution of signal energy across the different levels. After performing a full decomposition, one can calculate the variance (or energy, i.e., the sum of squared coefficients) of the detail coefficients at each level. High-frequency noise is typically characterized by low energy spread across the highest-frequency levels (e.g., D1, D2). The underlying signal and significant cyclical components, by contrast, will tend to have their energy concentrated in a few lower-frequency levels.[35] For this project, the initial hypothesis will be that the highest-frequency components, corresponding to 2-4 month cycles (Level 1) and 4-8 month cycles (Level 2), are dominated by noise and should be targeted for removal.

Critical Design Choice 3: Handling Boundary Effects

When applying a wavelet filter to a finite-length signal, the filter will inevitably extend beyond the start and end points of the data. This creates boundary or edge-effect artifacts, where the computed wavelet coefficients are distorted because they are influenced by non-existent data.[38] The region of the time-frequency plane affected by these artifacts is known as the "cone of influence".[40]

To mitigate this problem, the signal must be extended at its boundaries. The PyWavelets library provides several padding modes [12]:

- zero: Pads the signal with zeros, which can create an artificial jump at the boundary.
- constant: Pads with the last value, which can create an artificial plateau.
- periodic: Repeats the signal, which is rarely a valid assumption for financial time series.
- symmetric: Extends the signal by reflecting it at the boundary. This is often the most suitable choice for financial data as it typically creates the smoothest and most plausible extension, minimizing artificial discontinuities. Therefore, the analysis will use mode='sym'.

2.3 Step 3: Signal Denoising and Reconstruction

Once the signal is decomposed, the next step is to isolate the desired low-frequency component and reconstruct a "denoised" signal.

The Denoising and Reconstruction Process

The pywt.wavedec() function returns a list of coefficient arrays. The first element is the approximation coefficient array (cA) for the lowest-frequency level chosen, representing the long-term trend. The subsequent elements are the detail coefficient arrays (cD) for each level, ordered from the highest frequency (shortest period) to the lowest.[12]

The project proposes a simple and direct method for denoising:

1. **Decompose:** Apply the DWT to the momentum factor return series to a chosen level (e.g., level 4), yielding coefficients ``.
2. **Nullify High Frequencies:** Set the coefficients corresponding to the highest-frequency components to zero. For example, to remove the two highest-frequency levels, the coefficient list would be modified to ``.
3. **Reconstruct:** Use the inverse DWT function, pywt.waverec(), on the modified list of coefficients. The output is a new time series of the same length as the original, representing the "denoised" signal, which is essentially the original signal with its short-term oscillations filtered out.

An alternative, more sophisticated denoising technique involves *thresholding* the detail coefficients rather than simply setting them to zero.[12] In this approach, a threshold is calculated for each detail level. Coefficients with an absolute value below the threshold are considered noise and are set to zero (hard thresholding) or shrunk toward zero (soft thresholding), while coefficients above the threshold are retained. This method has the advantage of preserving large, significant short-term events (which would produce large detail coefficients) while still removing low-amplitude noise. While the core project will use the simpler nullification method, thresholding represents a valuable potential enhancement.

2.4 Step 4: Backtesting the Timing Strategy

The final step is to formulate and backtest a simple timing strategy based on the denoised signal and compare its performance against a passive benchmark.

Strategy Formulation

- **Benchmark Strategy:** A simple buy-and-hold investment in the constructed momentum factor portfolio. The portfolio is rebalanced monthly to maintain its long-short structure.
- **Wavelet-Timed Strategy:** This is a dynamic asset allocation strategy. At the end of each month t, the value of the reconstructed, denoised momentum signal is observed.
  - If the denoised signal for month t is **positive**, the strategy invests in the momentum factor portfolio for the upcoming month, t+1.
  - If the denoised signal for month t is **negative or zero**, the strategy exits the momentum factor and invests in a risk-free asset (e.g., 30-day US Treasury bills) for month t+1.

The intuition is that the denoised signal represents the underlying, persistent trend in momentum returns. A positive signal suggests the trend is upward and the factor is likely to perform well, while a negative signal suggests the trend is downward and a defensive position is warranted.

Performance Evaluation

To assess the efficacy of the timing strategy, a comprehensive set of performance metrics will be calculated for both the benchmark and the timed portfolios. The goal is not necessarily to maximize absolute returns, but to improve risk-adjusted performance.[41]

- **Key Performance Metrics:**
  - **Compound Annual Growth Rate (CAGR):** The geometric average annual rate of return.
  - **Annualized Volatility:** The standard deviation of monthly returns, scaled to an

annual figure.
  - **Sharpe Ratio:** The ratio of excess return (over the risk-free rate) to volatility. This is the primary measure of risk-adjusted return.
  - **Maximum Drawdown (MDD):** The largest peak-to-trough decline in portfolio value, representing the worst-case loss.
  - **Calmar Ratio:** The ratio of CAGR to MDD, measuring return relative to the worst-case loss.
  - **Monthly Turnover:** The percentage of the portfolio that is traded each month. This is crucial for assessing the potential impact of transaction costs.

The central research question is whether the wavelet-timed strategy can deliver a higher Sharpe Ratio and a lower Maximum Drawdown than the passive buy-and-hold benchmark. A successful outcome would demonstrate that the wavelet-based denoising process successfully extracts a valuable signal for tactical factor allocation.

Part III: Critical Analysis and Robustness Testing

A successful backtest, as outlined in Part II, is a necessary but insufficient condition for a viable quantitative strategy. The history of quantitative finance is littered with strategies that performed brilliantly in historical simulations but failed spectacularly in live trading. This failure often stems from two related issues: statistical biases in the research process and the omission of real-world market frictions. This section subjects the proposed wavelet timing strategy to a more rigorous critique, outlining the primary risks and proposing methods to test its robustness.

3.1 The Twin Perils: Overfitting and Data Snooping

Defining the Problem

- **Overfitting:** This occurs when a model is excessively complex relative to the amount of data available, causing it to "memorize" the random noise in the training data rather than learning the true underlying signal.[43] An overfit model will exhibit excellent performance on the historical data it was trained on (in-sample) but will fail to generalize and perform poorly on new, unseen data (out-of-sample).[45]
- **Data Snooping:** Also known as data dredging or p-hacking, this is the practice of conducting numerous statistical tests on a single dataset and selectively reporting only those results that appear statistically significant.[46] Given the probabilistic nature of statistical tests, if enough hypotheses are tested, some will appear significant purely by chance.[46] Data snooping is a primary driver of overfitting in quantitative finance, where

researchers may test thousands of potential strategies and report only the "winners".[48]

Sources of Overfitting in the Wavelet Timing Project

The proposed project, despite its apparent simplicity, contains numerous "researcher degrees of freedom"—parameters and methodological choices that can be tweaked. If a researcher iteratively adjusts these parameters to maximize the backtested Sharpe ratio, the final result is almost certain to be overfit. Key sources of this risk include:

1. **Factor Definition:** Minor changes to the momentum factor's construction (e.g., using a 12-month lookback vs. an 11-month lookback, changing the skip month, using different portfolio breakpoints) can alter the resulting time series.
2. **Choice of Mother Wavelet:** As shown in Table 2, there are many wavelet families (Daubechies, Symlets, Coiflets, etc.), each with numerous orders (e.g., db2, db4, db8). Testing all of them and picking the best performer is a classic case of data snooping.[35]
3. **Choice of Decomposition Level:** The number of levels to decompose the signal to, and which of those levels to discard as "noise," is a critical parameter. A researcher could test all possible combinations and select the one that yields the best backtest.[37]
4. **Denoising Method:** The choice between simple reconstruction and various forms of thresholding (soft, hard, different threshold selection rules) adds another layer of parameters to optimize.
5. **Timing Rule:** The binary rule (signal > 0) is just one possibility. One could test thresholds like signal > 0.001 or require the signal to be positive for two consecutive months, again creating an opportunity to fit the historical data.

Mitigation Strategy: A Rigorous, Hypothesis-Driven Approach

The most effective defense against data snooping is to adopt a disciplined, hypothesis-driven development process, akin to the scientific method.[50]

1. **A Priori Specification:** Before running a single backtest, the researcher must formally specify the complete strategy in a dedicated document. This includes the exact factor definition, the chosen mother wavelet and its justification (e.g., "a sym8 wavelet is chosen for its balance of smoothness and symmetry, which is expected to capture cyclical factor behavior without phase distortion"), the decomposition level (justified by theory, e.g., "levels 1 and 2, corresponding to cycles shorter than 8 months, are hypothesized to be noise"), and the precise timing rule.
2. **Single Test:** The backtest should then be run *once* on this pre-specified model. The result is a legitimate test of the initial hypothesis.
3. **Distinguish Confirmation from Exploration:** Any subsequent tests involving different parameters must be clearly labeled as *exploratory analysis*. The results of such exploration can be used to generate *new* hypotheses, but these new hypotheses must

then be validated on entirely different, unseen data (e.g., data from a different country or a future time period). Confusing exploratory findings with confirmatory results is the central error of data snooping.[47]

3.2 A More Realistic Backtest: Walk-Forward Optimization (WFO)

A single in-sample training period and out-of-sample testing period, while better than no out-of-sample testing at all, is often insufficient to establish robustness. The parameters optimized on the in-sample period may be uniquely suited to the specific market regimes present in that slice of history and may fail as market dynamics evolve.[51] Walk-Forward Optimization (WFO) is a more robust validation technique that simulates the real-world process of periodically re-evaluating and re-calibrating a model as new data becomes available.[52]

WFO Methodology

The WFO process involves a series of rolling backtests:

1. **Data Segmentation:** The entire historical dataset is divided into a number of contiguous "folds." For example, a 30-year dataset could be divided into 21 folds, each consisting of a 10-year "in-sample" training window followed by a 1-year "out-of-sample" testing window.
2. **Initial Optimization:** The strategy's parameters (e.g., the optimal wavelet decomposition level) are optimized on the first in-sample window (e.g., years 1-10) to maximize a target metric like the Sharpe ratio.
3. **First Out-of-Sample Test:** The "optimal" parameters found in step 2 are then applied, without change, to the first out-of-sample window (year 11). The performance during this period is recorded.
4. **Rolling Forward:** The window is then rolled forward by the length of the out-of-sample period. The new in-sample window becomes years 2-11. The parameters are re-optimized on this new window, and the resulting parameters are tested on the new out-of-sample window (year 12).
5. **Iteration:** This process is repeated until the end of the dataset is reached. The final performance of the strategy is the stitched-together time series of returns from all the individual out-of-sample periods.

This procedure tests the *robustness of the optimization process itself*. It answers the question: "If I had periodically re-calibrated my model based on the most recent decade of data, would it have consistently performed well in the subsequent year?" This is a much more realistic and stringent test than a simple static backtest.[54]

Pitfalls and Considerations for WFO

WFO is a powerful tool but not a silver bullet. It introduces its own set of design choices that can be sources of bias [51]:

- **Window Length:** The choice of the in-sample and out-of-sample window lengths is critical. A short in-sample window may not capture a full market cycle, leading to unstable parameters. A very long window may include outdated market regimes that are no longer relevant, biasing the optimization.[51]
- **Computational Cost:** WFO is computationally intensive, as it requires running multiple optimizations instead of just one.
- **Reactive Nature:** WFO adapts to market regime changes with a lag. Performance often deteriorates during a regime shift before the model can be re-optimized on data that includes the new regime.[51]

3.3 The Impact of Real-World Frictions

Even a robustly validated strategy can fail if its backtest ignores the unavoidable costs of trading. Factor timing strategies, by their nature, involve higher turnover than passive strategies, making them particularly sensitive to these costs.[55] A realistic backtest must incorporate models for these frictions.

Transaction Costs

Transaction costs can be broadly categorized into explicit and implicit costs.[57]

- **Explicit Costs:** These are direct, observable costs and include brokerage commissions, exchange fees, and taxes (e.g., stamp duty).[58] They are relatively straightforward to model, often as a fixed fee per trade or a percentage of the trade value.
- **Implicit Costs:** These are indirect and harder to measure but are often larger than explicit costs.
    - **Bid-Ask Spread:** The difference between the price at which one can sell an asset (the bid) and the price at which one can buy it (the ask). Any round-trip trade immediately incurs this cost.[59]
    - **Slippage:** This is the adverse price movement that occurs between the moment a trade decision is made and the moment the trade is actually executed. For a momentum strategy, which by definition buys assets whose prices are rising, slippage is a particularly significant cost, as the execution price is likely to be higher than the signal price.[58]
    - **Market Impact:** The effect that a large trade has on the market price. Executing a large buy order can push the price up, while a large sell order can depress it. This

cost is a function of the trade size relative to the asset's liquidity.[57]

Modeling Frictions in the Backtest

For the purposes of this project, a simple but effective transaction cost model can be implemented. After the WFO backtest generates a series of trades, the simulation can be re-run with the following adjustment: for every month that the strategy switches its position (e.g., from risk-free to momentum, or vice versa), a fixed cost in basis points (e.g., 5 bps or 0.05%) is subtracted from the monthly return. This single parameter can be used to approximate the combined effect of commissions, average spread, and slippage.
By running a sensitivity analysis—showing how the strategy's Sharpe ratio degrades as the assumed transaction cost increases from 0 bps to 10 bps, 20 bps, and beyond—one can determine the strategy's "breakeven" cost level. A strategy that is only profitable at unrealistically low cost assumptions (e.g., < 2 bps) is likely not viable in practice.
The progression from a naive backtest to a walk-forward analysis and finally to a friction-adjusted WFO establishes a clear hierarchy of evidence. A simple backtest provides the weakest evidence of a strategy's potential. A successful WFO demonstrates that the strategy's core logic is adaptable over time. A WFO that remains profitable after accounting for realistic trading costs provides the strongest evidence that the discovered effect may represent a genuine, exploitable market inefficiency rather than a statistical artifact.

Part IV: Advanced Concepts and Future Research

The core project outlined in the preceding sections represents a robust and insightful, yet foundational, application of wavelet analysis to factor timing. The true power of the wavelet toolkit, however, lies in its potential for more sophisticated applications that move beyond simple denoising. This section explores several advanced extensions, transforming the initial project from a standalone analysis into a launchpad for a comprehensive research agenda.

4.1 From Denoising to Regime Detection

The current project utilizes the wavelet decomposition for a single purpose: to filter out high-frequency noise and isolate a low-frequency trend. This process, however, discards a vast amount of potentially valuable information contained within the detail coefficients. The energy, or variance, of the wavelet coefficients at different scales provides a rich, time-varying signature of the signal's volatility structure, which can be used to identify market regimes.[13]
For instance, a period where the energy is concentrated in the high-frequency detail coefficients (D1, D2) corresponds to a choppy, noisy market regime, whereas a period where energy is concentrated in lower-frequency coefficients (D4, D5) indicates a smoother,

trending regime.

Proposed Extension: Wavelet-Based Regime-Conditioned Timing

A more advanced strategy could be developed by using the full wavelet power spectrum to classify the market state.

1. **Regime Classification:** At each point in time, calculate the relative energy of the wavelet coefficients at each decomposition level over a rolling window. Use these energy distributions as features in a clustering algorithm (e.g., k-means) to classify the current market environment into a set of discrete regimes (e.g., "Low-Volatility Trend," "High-Volatility Mean-Reversion," "Noisy/Choppy").
2. **Dynamic Strategy Switching:** The timing strategy can then be made regime-dependent. For example:
    - In a "Low-Volatility Trend" regime, apply the original timing rule based on the denoised signal's direction.
    - In a "High-Volatility Mean-Reversion" regime, switch to a contrarian strategy or reduce exposure entirely.
    - In a "Noisy/Choppy" regime, remain in a risk-off position, as timing signals are likely to be unreliable.

This approach aligns with cutting-edge research that uses wavelet-based representations to classify distinct types of market events. For example, studies have shown that the time-asymmetry of volatility around a price jump, a feature readily captured by wavelets, can distinguish between endogenous jumps (caused by internal market dynamics) and exogenous jumps (caused by external news).[60] A highly sophisticated timing model could be designed to filter its signals based on the detected nature of recent market shocks, for instance, by being more aggressive after exogenous shocks and more cautious during periods of endogenous instability.

4.2 Next-Generation Signal Extraction: Element Analysis

While standard DWT-based denoising is effective, it has inherent limitations. Noise components can still be mapped to wavelet scales, and true signal features that do not perfectly align with the shape of the mother wavelet can be "smeared" across multiple decomposition levels, resulting in a blurred and imperfectly reconstructed signal.[10]

A Superior Method: Element Analysis

A novel technique known as "Element Analysis" offers a more powerful approach to signal extraction in noisy environments.[63] Instead of modeling a time series as a sum of scaled and shifted wavelets, Element Analysis models it as a baseline of Gaussian noise upon which discrete, isolated "events" or "generators" are superimposed. These events produce ripples of various frequencies. The method's goal is not to perfectly reconstruct the entire signal, but to directly estimate the parameters (time, scale, phase, amplitude) of these underlying generating functions, while explicitly rejecting perturbations that are statistically indistinguishable from noise.[10]

Proposed Extension: Timing with an Element-Analyzed Signal

The core project could be re-executed using Element Analysis as the signal extraction engine instead of DWT denoising.

1. **Model the Factor Series:** Model the momentum factor return series as a collection of significant "events" over a noise floor using Element Analysis.
2. **Reconstruct the "Generator" Signal:** Reconstruct a time series based only on the parameters of the identified event generators.
3. **Backtest the Timing Strategy:** Use this new, highly purified signal as the input for the same positive/negative timing rule.

The central hypothesis is that this method will produce a cleaner and more robust signal, as it is specifically designed to distinguish true, localized signal events from the pervasive noise characteristic of financial data. Empirical applications have shown that Element Analysis can identify significant volatility events that are completely obscured by noise in traditional wavelet scalograms.[63] This suggests it could provide a superior foundation for a timing model.

4.3 Multi-Factor and Cross-Asset Applications

The project's focus on a single factor, momentum, is a necessary simplification for a two-week study. However, the true power of time-frequency analysis may be most evident in multivariate contexts, where it can be used to understand the complex, time-varying relationships between different assets or factors.[7]

Wavelet Coherence

Wavelet coherence is a tool that extends the concept of correlation into the time-frequency domain. It measures the cross-correlation between two time series as a function of both time and frequency (or scale). A wavelet coherence plot can reveal, for example, that two factors are strongly and positively correlated during short-term market panics (a high-frequency relationship localized in time) but are negatively correlated in their long-term trends (a persistent low-frequency relationship).[65]

Proposed Extension: A Dynamic Multi-Factor Allocation Model

1. **Denoise Multiple Factors:** Apply the wavelet denoising technique developed in the core project to a suite of standard factors, such as Value, Quality, and Low Volatility, in addition to Momentum.
2. **Analyze Time-Scale Relationships:** Use wavelet coherence to analyze the pair-wise relationships between the *denoised* long-term components of these factors. This analysis would seek to answer questions like:
   - Does the well-known negative long-term correlation between Value and Momentum hold consistently across time, or does it break down in certain regimes?
   - How does the relationship between defensive factors (Low Volatility, Quality) and cyclical factors (Value, Momentum) change across different time scales?
3. **Build a Dynamic Allocation Model:** The insights from the coherence analysis could inform a sophisticated factor allocation model. Instead of simply timing individual factors in isolation, the model could dynamically adjust allocations based on the prevailing time-scale correlation structure. For example, if the long-term trends of Value and Momentum become positively correlated (a rare but possible regime), the model might reduce exposure to both to avoid unintended concentration of risk.

This progression—from simple denoising of a single series, to regime-aware timing, to superior signal extraction, and finally to multivariate dynamic allocation—forms a logical and compelling research roadmap. It demonstrates that the initial project is not an isolated exercise but a foundational step into a rich and promising area of modern quantitative finance.

Conclusion

This report has provided a comprehensive blueprint for a two-week quantitative research project centered on using wavelet analysis to time the momentum factor. The investigation proceeds from a firm theoretical grounding to a detailed practical implementation plan, a critical analysis of statistical pitfalls, and an exploration of advanced research avenues.

The theoretical justification for using wavelets is compelling. Financial time series are characterized by non-stationarity and transient events, features that are poorly handled by traditional methods like Fourier analysis or the Hodrick-Prescott filter. Wavelet analysis, with its inherent ability to localize signals in both time and frequency, offers a flexible and data-driven framework for decomposing a factor return series into its underlying trend and higher-frequency noise components.

The practical implementation, however, is laden with critical methodological choices. The construction of the momentum factor itself requires a meticulous process of data acquisition, linking, and cleaning, where each decision can materially impact the final time series. The application of the wavelet transform necessitates principled choices regarding the mother wavelet, the decomposition level, and the handling of boundary effects. A naive backtest of a simple timing strategy based on the denoised signal might show promising results, but such a finding must be treated with extreme skepticism.

The most critical phase of the project involves rigorous robustness testing. The numerous "researcher degrees of freedom" in the analysis create a significant risk of overfitting and data snooping. A disciplined, hypothesis-driven approach, coupled with a more realistic backtesting methodology like Walk-Forward Optimization, is essential to validate the strategy's adaptability to changing market conditions. Furthermore, the inclusion of realistic transaction costs is non-negotiable, as the higher turnover of timing strategies can easily erode or eliminate theoretical profits.

The final performance of the strategy, after these rigorous tests, provides a more honest assessment of its potential. The table below summarizes the hypothetical performance across different stages of validation, illustrating the typical degradation of performance as more realistic assumptions are introduced.

Table 3: Backtest Performance Summary (Hypothetical Results)

Performance Metric	Momentum Factor (Benchmark)	Wavelet-Timed (Static, No Costs)	Wavelet-Timed (WFO, No Costs)	Wavelet-Timed (WFO, With Costs)
CAGR	8.5%	7.8%	7.2%	6.1%
Annualized Volatility	16.0%	11.5%	12.0%	12.1%
Sharpe Ratio	0.53	0.68	0.60	0.50
Maximum Drawdown	-55.0%	-25.0%	-28.0%	-29.5%
Monthly Turnover	15.0%	8.0%	9.5%	9.5%

As illustrated, while the initial, naive timing strategy might show a significant improvement in the Sharpe ratio and a dramatic reduction in drawdown, its advantage shrinks under the more robust WFO framework and may even fall below the benchmark after accounting for transaction costs. This outcome is common in quantitative research and underscores the central message: **the process of validation is more important than the initial result.**

Ultimately, this project serves as an excellent case study in modern quantitative research. It demonstrates how techniques from other scientific disciplines can provide novel insights into financial markets. More importantly, it highlights that the path from a creative idea to a viable investment strategy is one of disciplined hypothesis testing, rigorous defense against statistical biases, and a healthy respect for the frictions of the real world. While the simple timing strategy proposed here may not be a "holy grail," the analytical framework it introduces—moving from denoising to regime detection and multivariate analysis—opens the door to a rich and potentially fruitful domain of advanced factor modeling.

Works cited

1. Introduction to Non-Stationary Processes - Investopedia, accessed July 30, 2025, https://www.investopedia.com/articles/trading/07/stationary.asp
2. Understanding the Importance of Stationarity in Time Series - Hex, accessed July 30, 2025, https://hex.tech/blog/stationarity-in-time-series/
3. Wavelet Analysis and Applications in Economics and Finance | Open Access Journals, accessed July 30, 2025, https://www.rroij.com/open-access/wavelet-analysis-and-applications-in-economics-and-finance.php?aid=62323
4. Denoise of Financial Time Series Data using Wavelet Transform [closed] - Stack Overflow, accessed July 30, 2025, https://stackoverflow.com/questions/54851568/denoise-of-financial-time-series-data-using-wavelet-transform
5. An Application of the Continuous Wavelet Transform to Financial Time Series - CORE, accessed July 30, 2025, https://core.ac.uk/download/pdf/289954260.pdf
6. TimeBridge: Non-Stationarity Matters for Long-term Time Series Forecasting - arXiv, accessed July 30, 2025, https://arxiv.org/html/2410.04442v1
7. Wavelet Multiresolution Analysis of Financial Time Series - CORE, accessed July 30, 2025, https://core.ac.uk/download/pdf/197968234.pdf
8. Wavelets in Economics - Banco de Portugal, accessed July 30, 2025, https://www.bportugal.pt/sites/default/files/anexos/papers/ab201208_e.pdf
9. Wavelet analysis and synthesis of fractional Brownian motion - Information Theory, IEEE Transactions on, accessed July 30, 2025, https://perso.ens-lyon.fr/patrick.flandrin/IEEE_IT1992.pdf
10. Wavelet Analysis for Time Series Financial Signals via Element Analysis - ResearchGate, accessed July 30, 2025, https://www.researchgate.net/publication/367652258_Wavelet_Analysis_for_Time_Series_Financial_Signals_via_Element_Analysis
11. Wavelet Analysis of Financial Data - MATLAB & Simulink Example - MathWorks, accessed July 30, 2025, https://www.mathworks.com/help/wavelet/ug/wavelet-analysis-of-financial-data.html
12. Practical Wavelet Techniques for Time Series - Number Analytics, accessed July 30, 2025, https://www.numberanalytics.com/blog/practical-wavelet-techniques-time-series
13. Wavelet Analysis and Active Media Technology : DETECT MARKET ..., accessed July 30, 2025, https://www.worldscientific.com/doi/pdf/10.1142/9789812701695_0231
14. Boosting: Why you Can Use the HP Filter, accessed July 30, 2025, https://cowles.yale.edu/sites/default/files/2022-08/d2212.pdf
15. Hodrick–Prescott filter - Wikipedia, accessed July 30, 2025, https://en.wikipedia.org/wiki/Hodrick%E2%80%93Prescott_filter
16. The Hodrick-Prescott Filter at Time Series Endpoints - University of Nottingham, accessed July 30, 2025, https://www.nottingham.ac.uk/economics/documents/discussion-papers/03-08.p

df
17. Kalman filter - Wikipedia, accessed July 30, 2025, https://en.wikipedia.org/wiki/Kalman_filter
18. Kalman Filter and Time Series, accessed July 30, 2025, https://digilib.uhk.cz/bitstream/handle/20.500.12603/102/FRON%C4%8CKOV%C3%81%2C%20PRA%C5%BD%C3%81K.pdf?sequence=1&isAllowed=y
19. Predicting Time SeriesUsing an Automatic New Algorithm of the Kalman Filter - MDPI, accessed July 30, 2025, https://www.mdpi.com/2227-7390/10/16/2915
20. On the Construction of Common Size, Value and Momentum Factors in International Stock Markets: A Guide with Applications - EconStor, accessed July 30, 2025, https://www.econstor.eu/bitstream/10419/171584/1/wp-11-141.pdf
21. Momentum and Value Reimagined - CBS Research Portal, accessed July 30, 2025, https://research.cbs.dk/files/108049160/1850842_Thesis.pdf
22. Global Factor Data Documentation - AWS, accessed July 30, 2025, https://jkpfactors.s3.amazonaws.com/documents/Documentation.pdf
23. CRSP/Compustat Merged Database - WRDS, accessed July 30, 2025, https://wrds-www.wharton.upenn.edu/documents/398/CRSP_-_Compustat_Merged_Database_CCM.pdf
24. Data Quality Problems Troubling Business and Financial Researchers: A Literature Review and Synthetic Analysis - Digital Commons @ West Chester University, accessed July 30, 2025, https://digitalcommons.wcupa.edu/cgi/viewcontent.cgi?article=1013&context=lib_facpub
25. Does the source of fundamental data matter? - EconStor, accessed July 30, 2025, https://www.econstor.eu/bitstream/10419/203194/1/1027882285.pdf
26. Lost in Standardization: Revisiting Accounting-based Return ..., accessed July 30, 2025, https://business.columbia.edu/sites/default/files-efs/imce-uploads/CEASA/Events%20Page/revisiting_accounting-based_return_anomalies.pdf
27. Does Compustat Data Standardization Improve Bankruptcy Prediction Models? | Request PDF - ResearchGate, accessed July 30, 2025, https://www.researchgate.net/publication/272241981_Does_Compustat_Data_Standardization_Improve_Bankruptcy_Prediction_Models
28. Compustat vs Factset Factor Values Don't Match - Portfolio123 Community, accessed July 30, 2025, https://community.portfolio123.com/t/compustat-vs-factset-factor-values-dont-match/58667
29. Compustat Data: A Misleading Measure of Corporate Market Power and Market Competition, accessed July 30, 2025, https://itif.org/publications/2024/08/12/compustat-data-misleading-measure-corporate-market-competition/
30. The strategic value of data hygiene: Guide for finance leaders - Cube Software, accessed July 30, 2025, https://www.cubesoftware.com/blog/data-hygiene-for-finance
31. Chapter 5 Cleaning and processing research data | Development Research in

Practice The DIME Analytics Data Handbook - worldbank.github.io, accessed July 30, 2025, https://worldbank.github.io/dime-data-handbook/processing.html
32. The Wharton Research Data Service (WRDS) data set and Factor Model Factors, accessed July 30, 2025, http://bearcave.com/finance/wrds_data.html
33. Kenneth R. French - Detail for Monthly Momentum Factor (Mom), accessed July 30, 2025, https://mba.tuck.dartmouth.edu/pages/faculty/ken.french/Data_Library/det_mom_factor.html
34. Data Cleaning: Definition, Benefits, And How-To | Tableau, accessed July 30, 2025, https://www.tableau.com/learn/articles/what-is-data-cleaning
35. (PDF) A comprehensive guide to selecting suitable wavelet decomposition level and functions in discrete wavelet transform for fault detection in distribution networks - ResearchGate, accessed July 30, 2025, https://www.researchgate.net/publication/387798284_A_comprehensive_guide_to_selecting_suitable_wavelet_decomposition_level_and_functions_in_discrete_wavelet_transform_for_fault_detection_in_distribution_networks
36. Can Denoising Enhance Prediction Accuracy of Learning Models? A Case of Wavelet Decomposition Approach - MDPI, accessed July 30, 2025, https://www.mdpi.com/2571-9394/6/1/5
37. Discussion on the Choice of Decomposition Level for Wavelet Based Hydrological Time Series Modeling - MDPI, accessed July 30, 2025, https://www.mdpi.com/2073-4441/8/5/197
38. Boundary Effects Reduction in Wavelet Transform for Time-frequency Analysis - WSEAS, accessed July 30, 2025, https://www.wseas.org/multimedia/journals/signal/2012/53-661.pdf
39. Boundary effects reduction in wavelet transform for time-frequency analysis - ResearchGate, accessed July 30, 2025, https://www.researchgate.net/publication/290026178_Boundary_effects_reduction_in_wavelet_transform_for_time-frequency_analysis
40. Boundary Effects and the Cone of Influence - MATLAB & Simulink - MathWorks, accessed July 30, 2025, https://www.mathworks.com/help/wavelet/ug/boundary-effects-and-the-cone-of-influence.html
41. Avoid Equity Bear Markets with a Market Timing Strategy – Part 1 - QuantPedia, accessed July 30, 2025, https://quantpedia.com/avoid-equity-bear-markets-with-a-market-timing-strategy-part-1/
42. Market Timing Strategies | (Setups, Regime Filters & Backtest) - QuantifiedStrategies.com, accessed July 30, 2025, https://www.quantifiedstrategies.com/market-timing/
43. Overfitting, Model Tuning, and Evaluation of Prediction Performance - Multivariate Statistical Machine Learning Methods for Genomic Prediction - NCBI, accessed July 30, 2025, https://www.ncbi.nlm.nih.gov/books/NBK583970/
44. Overfitting | Machine Learning | Google for Developers, accessed July 30, 2025, https://developers.google.com/machine-learning/crash-course/overfitting/overfitt

ing
45. What is Overfitting? - Overfitting in Machine Learning Explained - AWS, accessed July 30, 2025, https://aws.amazon.com/what-is/overfitting/
46. Data snooping - Stanford Data Science, accessed July 30, 2025, https://datascience.stanford.edu/news/data-snooping
47. Understanding Data Snooping: Key Techniques to Prevent Analysis Bias - Number Analytics, accessed July 30, 2025, https://www.numberanalytics.com/blog/understanding-data-snooping-techniques-prevent-analysis-bias
48. Data Snooping Bias - Meaning, Examples, How to Avoid? - WallStreetMojo, accessed July 30, 2025, https://www.wallstreetmojo.com/data-snooping-bias/
49. We Found 30 Timing Strategies that "Worked"—and 690 that Didn't | Dimensional, accessed July 30, 2025, https://www.dimensional.com/us-en/insights/we-found-30-timing-strategies-that-worked-and-690-that-didnt
50. (PDF) Developing & Backtesting Systematic Trading Strategies, accessed July 30, 2025, https://www.researchgate.net/publication/319298448_Developing_Backtesting_Systematic_Trading_Strategies
51. Walk-Forward Optimization: How It Works, Its Limitations, and Backtesting Implementation, accessed July 30, 2025, https://blog.quantinsti.com/walk-forward-optimization-introduction/
52. How to Use Walk Forward Analysis: You May Be Doing It Wrong! - Unger Academy EN, accessed July 30, 2025, https://ungeracademy.com/posts/how-to-use-walk-forward-analysis-you-may-be-doing-it-wrong
53. Mastering Walk-Forward Optimization - Number Analytics, accessed July 30, 2025, https://www.numberanalytics.com/blog/mastering-walk-forward-optimization
54. Refining and Robust Backtesting of A Century of Profitable Industry Trends - arXiv, accessed July 30, 2025, https://arxiv.org/html/2412.14361v2
55. Factor Timing in Asset Management: A Literature Review - ResearchGate, accessed July 30, 2025, https://www.researchgate.net/publication/388107691_Factor_Timing_in_Asset_Management_A_Literature_Review
56. A prudent route to effective factor timing - Robeco.com, accessed July 30, 2025, https://www.robeco.com/files/docm/docu-202307-a-prudent-route-to-effective-factor-timing.pdf
57. Transaction Cost Modeling - QuestDB, accessed July 30, 2025, https://questdb.com/glossary/transaction-cost-modeling/
58. Successful Backtesting of Algorithmic Trading Strategies - Part II | QuantStart, accessed July 30, 2025, https://www.quantstart.com/articles/Successful-Backtesting-of-Algorithmic-Trading-Strategies-Part-II/
59. How to Avoid Common Mistakes in Backtesting? - Quantra by QuantInsti,

accessed July 30, 2025, https://quantra.quantinsti.com/glossary/How-to-Avoid-Common-Mistakes-in-Backtesting
60. Identifying new classes of financial price jumps with wavelets - PNAS, accessed July 30, 2025, https://www.pnas.org/doi/10.1073/pnas.2409156121
61. Riding Wavelets: A Method to Discover New Classes of Price Jumps - arXiv, accessed July 30, 2025, https://arxiv.org/html/2404.16467v1
62. Wavelet Analysis for Time Series Financial Signals via Element Analysis - arXiv, accessed July 30, 2025, https://arxiv.org/pdf/2301.13255
63. Wavelet Analysis for Time Series Financial Signals via Element ..., accessed July 30, 2025, https://arxiv.org/abs/2301.13255
64. TIME–FREQUENCY ANALYSIS BETWEEN ECONOMIC RISK AND, accessed July 30, 2025, https://worldscientific.com/doi/10.1142/S2010495222500130
65. Wavelet analysis and energy-based measures for oil-food price relationship as a footprint of financialisation effect - arXiv, accessed July 30, 2025, https://arxiv.org/pdf/2104.11891